Summary: Inflation Attention Thresholds and Regional Phillips Curves

Zorah Zafari

June 2025

Key Results

This paper empirically tests the inflation attention threshold hypothesis using household-level panel data from the University of Michigan Survey of Consumers, merged with regional CPI inflation data. The central hypothesis is that household inflation expectation revisions become more pronounced when local inflation exceeds a 4% threshold.

The baseline regression confirms this hypothesis: households in high-inflation regions revise their short-term inflation expectations more sharply. This result is statistically significant even after controlling for regional and time fixed effects and a suite of demographic covariates, such as age, education, and income.

When disaggregating by education level, the threshold effect is strongest among those with some high school education. College-educated respondents do not show significant responsiveness to high-inflation regimes, suggesting that more educated individuals form more stable expectations or are less sensitive to local price changes.

Similarly, income-based heterogeneity reveals that the effect is concentrated among households in the bottom income quartiles. The top 25% income group shows no significant revision behavior when inflation exceeds 4%, implying that economic vulnerability may amplify sensitivity to salient inflation signals.

Conceptual and Methodological Contributions

1. What are we learning at the household level that Pfäuti (2025) does not show at the aggregate?

While Pfäuti (2025) demonstrates the existence of an inflation attention threshold using aggregated version of the same survey, my paper exploits the micro-level panel structure by linking respondents

across interviews. The use of panel data allows me to track the same household over time, controlling for fixed unobservables, and showing that within-household expectation revisions respond differently depending on local inflation exposure. This allows me to directly observe heterogeneity in attention behavior across income and education groups—something aggregate regressions obscure.

2. What particular confounder am I controlling for—i.e., state dependence?
State dependence is a key concern: it may be that households revise expectations more under high inflation simply because these regimes are more volatile. By including region and time fixed effects, and conditioning on detailed demographics, I control for unobserved persistent characteristics that could confound the threshold effect. My approach ensures that identification comes from variation in inflation exposure and expectation updates within region-time cells.

3. Why is this the right regression to run?
This regression directly tests the hypothesis that expectation updating is nonlinear and intensifies above a threshold. The dependent variable—absolute revision in expected inflation—measures attention directly. The main regressor, a binary high-inflation indicator, follows Pfäuti's theoretical cutoff. Region and time fixed effects address macro confounds. Subgroup regressions then test for heterogeneity. This clean and interpretable setup balances internal validity with external relevance.

4. What am I doing differently than Oliver Pfäuti?
Pfäuti provides compelling macro-level correlations and theoretical motivation. I empirically test the micro-foundations he posits—linking regional inflation to individual behavior—and extend his insights by showing which types of households drive threshold-based attention effects. I also incorporate regional heterogeneity, which Pfäuti acknowledges but does not exploit directly.

Conclusion and Future Directions

This paper provides the first micro-level validation of the inflation attention threshold hypothesis. Using panel data, I show that household inflation expectation revisions are significantly larger when local inflation exceeds 4%. The results are strongest for low-income and less-educated households, supporting the idea that attention is both state- and agent-dependent.

Looking ahead, future work will refine the threshold itself—testing for optimal or group-specific cutoffs using nonparametric and regime-switching approaches. I also plan to model expectation revision as a function of lagged inflation volatility and use interaction terms to formally test for state dependence. These extensions will enhance our understanding of attention formation and improve macroeconomic models that rely on expectations as a core transmission channel.